CROP YIELD PREDICTION USING MACHINE LEARNING

A CAPSTONE PROJECT REPORT

Submitted in partial fulfillment of the
requirement for the award of the
Degree of

BACHELOR OF TECHNOLOGY
IN
COMPUTER SCIENCE & ENGINEERING

by

J. ARTHI PRASANNA (21BCE9721)
K. HARSHA VARDHAN NAIK (21BCE9743)
M. SATISH (21BCE9633)
B. SATWIK (21BCE9305)

Under the Guidance of

DR. RAM MOHAN

SCHOOL OF COMPUTER SCIENCE AND ENGINEERING
VIT-AP UNIVERSITY
AMARAVATI- 522237

DECEMBER 2024

CERTIFICATE

This is to certify that the Capstone Project work titled "**CROP YIELD PREDICTION USING MACHINE LEARNING**" that is being submitted by **J. ARTHI PRASANNA (21BCE9721), K. HARSHA VARDHAN NAIK (21BCE9743), M. SATISH (21BCE9633), and B. SATWIK (21BCE9305)** is in partial fulfillment of the requirements for the award of Bachelor of Technology, is a record of bonafide work done under my guidance. The contents of this Project work, in full or in parts, have neither been taken from any other source nor have been submitted to any other Institute or University for award of any degree or diploma and the same is certified.

Dr. Ram Mohan
Guide

The thesis is satisfactory / unsatisfactory

Posham Uppamma
Internal Examiner1

Priyanka Singh
Internal Examiner2

Approved

HoD, Department of Artificial Intelligence and Machine Learning,
School of Computer Science and Engineering

ACKNOWLEDGEMENTS

We express our gratitude to Dr.Ram Mohan our mentor, for his unwavering encouragement and assistance during the whole endeavour. He would frequently ask me questions about specific parts of the project that would help me figure out why and how to solve particular problems. His unwavering support has had a great effect on me, and his concepts are reflected in the things I do. We also want to thank the School of Computer Science and Engineering for giving me the chance to improve my abilities and expand my skill set, both of which will undoubtedly help me in the years to come. A project serves as a scaffold between theoretical and practical learning. With this in mind, we worked hard and made progress on the endeavour, with the support and encouragement of everyone in our immediate vicinity. Yes, I would also want to thank my supervisors and friends for their encouragement and support in helping me organize and come up with creative solutions for my assignments. I am really appreciative of each of these. We were able to complete this work and make it a wonderful and enjoyable experience because of them. This project is the result of many people's tireless labour, without which it would not have been feasible

ABSTRACT

Crop yield prediction plays a crucial role in modern agriculture, helping farmers and policymakers optimize resources and make informed decisions. This study explores the application of machine learning techniques to predict crop yields accurately based on environmental, soil, and historical agricultural data. By leveraging algorithms such as Random Forest, Support Vector Machines, and Neural Networks, the model analyzes factors like rainfall, temperature, soil type, and farming practices to provide reliable yield forecasts. The proposed approach aims to enhance agricultural efficiency, mitigate risks, and contribute to food security. Experimental results demonstrate the model's effectiveness, offering a significant improvement in prediction accuracy compared to traditional methods. This work underscores the potential of data-driven solutions in addressing challenges in agriculture and supporting sustainable practices.

TABLE OF CONTENTS

List of Figures

CHAPTER 1

INTRODUCTION

Crop yield prediction is a critical aspect of agricultural planning, enabling farmers and policymakers to make informed decisions about resource allocation, crop selection, and risk management. Traditional methods of predicting crop yields rely heavily on manual analysis, historical trends, and localized expertise, which are often inadequate in addressing the complexities of modern agriculture. Factors such as climate change, fluctuating weather patterns, and soil variability add layers of uncertainty, making accurate predictions challenging.

Machine learning (ML) offers a data-driven approach to tackle these challenges by leveraging large datasets and advanced computational techniques. ML algorithms can analyze diverse factors such as climatic conditions, soil properties, crop varieties, and historical yield data to identify patterns and relationships that are not easily discernible through traditional methods. Among various ML techniques, Decision Tree algorithms stand out for their simplicity, interpretability, and ability to handle non-linear relationships.

In this project, we explore the use of a Decision Tree-based machine learning model for crop yield prediction. The model incorporates multiple parameters, including rainfall, temperature, soil pH, and historical crop yields, to generate reliable and actionable predictions. The integration of machine learning into agricultural practices has the potential to enhance productivity, reduce resource wastage, and contribute to sustainable farming, ultimately addressing the growing demand for food in a changing global environment.

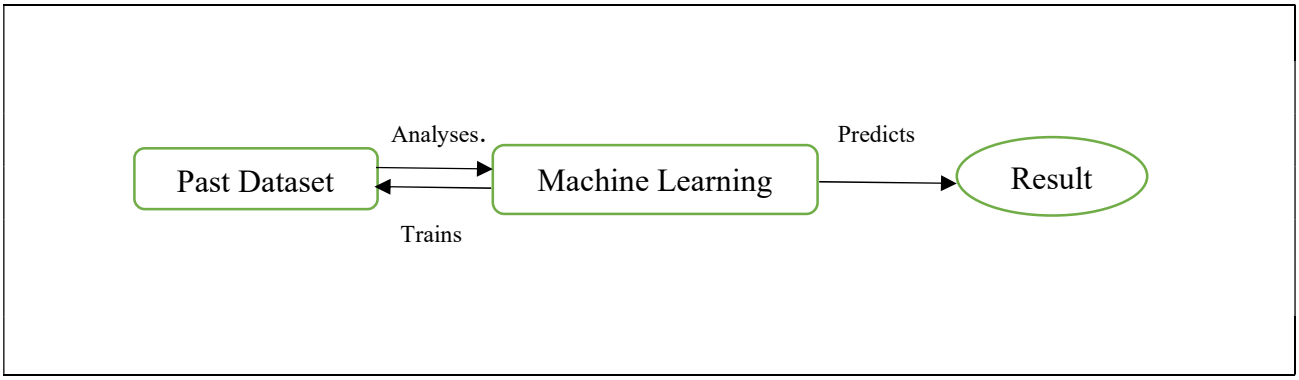

Fig 1.1 Process of Machine Learning

1.1 Objectives

The following are the objectives of this project:

- **Develop an Accurate Prediction Model**: Utilize machine learning algorithms, particularly Decision Trees, to predict crop yields with high accuracy.
- **Incorporate Multi-Factor Analysis**: Analyze diverse factors such as climate conditions, soil properties, and historical yield data to improve prediction reliability.
- **Support Resource Optimization**: Enable farmers to optimize the use of water, fertilizers, and other resources based on predicted yields.
- **Enhance Agricultural Planning**: Assist policymakers and stakeholders in making informed decisions regarding crop selection, planting schedules, and land use.
- **Mitigate Risks**: Reduce the uncertainty associated with weather variability and climate change by providing data-driven insights.
- **Promote Sustainable Farming**: Encourage practices that align with sustainable agriculture by minimizing resource wastage and improving productivity.
- **Facilitate Scalability**: Develop a model that can be scaled and adapted to different regions and crops, making it widely applicable.
- **Enable Real-Time Decision Support**: Provide predictions that can be used in real-time to guide day-to-day agricultural activities.

1.2 Background and Literature Survey

Crop yield prediction has been a critical area of research in agriculture, aiming to address challenges such as food security, efficient resource allocation, and agricultural sustainability. Traditionally, yield predictions were made using statistical methods and expert knowledge, which often failed to capture the complex, non-linear interactions among climatic, soil, and agricultural factors. The advent of machine learning (ML) has transformed this domain, offering sophisticated tools to analyze large datasets and uncover patterns that can improve prediction accuracy. Machine learning techniques have enabled the development of data-driven models that integrate multiple parameters, such as historical crop yields, weather conditions, soil characteristics, and agricultural practices. Among these, Decision Tree algorithms are particularly notable for their ability to model complex relationships while remaining interpretable. This makes them suitable for use in agriculture, where end-users, such as farmers and policymakers, require transparent and actionable insights.

Literature Survey

Earlier studies predominantly used statistical methods, such as linear regression and time series analysis, for yield prediction. While effective for small datasets, these methods often underperformed in handling large-scale and multi-dimensional agricultural data.

Machine Learning in Agriculture:

- *Random Forest and Gradient Boosting*: Research has shown these ensemble methods to be effective in crop yield prediction due to their robustness and high accuracy. However, they often lack interpretability, which limits their usability for non-technical stakeholders.
- *Support Vector Machines*: These have been employed for yield prediction, particularly in scenarios where data dimensions are high. However, their computational complexity can be a challenge.

Studies have highlighted the advantages of Decision Trees in agricultural applications, emphasizing their simplicity, transparency, and ability to handle categorical and continuous data.

- **Niazian et al. (2019)**: Demonstrated the use of Decision Trees for predicting wheat yield based on environmental and soil variables.
- **Patel et al. (2021)**: Applied Decision Tree algorithms to predict rice yields, showing improved results compared to traditional statistical methods.

Hybrid

Recent studies combine Decision Trees with ensemble methods (e.g., Random Forests) or neural networks to achieve higher accuracy while retaining interpretability. For example, integrating Decision Trees with fuzzy logic has been explored for yield prediction under uncertain climatic conditions.

Real-World

Applications of ML models for yield prediction have been successfully implemented in various countries, such as India, the USA, and Brazil, focusing on crops like rice, wheat, and maize. These models often integrate satellite imagery, IoT data, and weather forecasts to enhance their predictions.

1.3 Organization of the Report

The remaining chapters of the project report are described as follows:

- Chapter 2 contains the proposed system, methodology, hardware and software details, system architecture.
- Chapter 3 project requirements, implementation of the project, system flow design.
- Chapter 4 discusses the results obtained after the project was implemented.
- Chapter 5 conclusion and future work
- Chapter 6 consists of codes.
- Chapter 7 gives references.

CHAPTER 2

CROP YIELD PREDICTION USING MACHINE LEARNING

This Chapter describes the proposed system, working methodology, software and hardware details.

2.1 Proposed System

The system is designed to predict crop yields based on environmental and agricultural inputs using a Decision Tree machine learning model.

1. Data Preprocessing

- **Handling Missing Values**: Ensure no missing values in critical features such as rainfall, temperature, and yield.
- **Encoding Categorical Variables**: Convert categorical data, like "Area" and "Item," into numerical formats using techniques such as one-hot encoding or label encoding.
- **Feature Scaling**: Normalize numerical features to ensure consistent scaling for model training.

2. Feature Selection

- Identify the most relevant features (e.g., rainfall, temperature, and pesticide use) for predicting crop yields using statistical techniques and correlation analysis.

3. Model Development

- **Algorithm:** Implement a Decision Tree regression model to predict crop yields.
- **Training:** Train the model on historical data using features such as "average_rain_fall_mm_per_year," "pesticides_tonnes," and "avg_temp."
- **Testing and Validation:** Evaluate the model using metrics like Mean Absolute Error (MAE) and R-squared to assess accuracy.

4. a. Input: Collect inputs for rainfall, temperature, pesticide use, and crop type for a given area and year.

b. Processing: The system processes the input data and uses the trained Decision Tree model to predict crop yields.

c. Output: Provide predicted yield values for the given inputs.

5. Deployment

- Develop an interactive user interface (e.g., web or mobile app) where users can input parameters and receive yield predictions.
- Incorporate a database to store historical and predicted data for future analysis.

6. Benefits

- Helps farmers optimize resources by predicting yields based on environmental conditions.
- Assists policymakers in planning and resource allocation.
- Encourages sustainable agricultural practices by providing data-driven insights.

2.2 Methodology

The methodology involves several structured steps to build and implement a machine learning model for predicting crop yields.

1. Problem Understanding

- Define the objectives of crop yield prediction (e.g., improving accuracy, aiding decision-making).
- Identify key factors influencing yield (e.g., climatic, soil, and agricultural inputs).

2. Data Collection and Understanding

- Utilize the provided dataset, which includes features such as average rainfall, average temperature, pesticide usage, crop type, and yield.
- Explore the dataset to understand its structure, identify patterns, and detect anomalies.

3. Data Preprocessing

- **Data Cleaning**:
  - Handle missing or inconsistent values in key features.
  - Remove duplicates and irrelevant records.
- **Feature Encoding:**
  - Convert categorical variables like "Area" and "Item" into numerical formats using label encoding or one-hot encoding.
- **Feature Scaling**:
  - Normalize numerical features (e.g., rainfall, temperature, and pesticide usage) to ensure consistency in data distribution.
- **Feature Engineering**:
  - Derive new features if needed, such as rainfall deviation from the average or combined effect of temperature and rainfall.

4. Exploratory Data Analysis (EDA)

- Visualize relationships between features and crop yield using scatter plots, heatmaps, and boxplots.
- Identify correlations among variables to select relevant features**.**

5. Model Development

- Algorithm Selection: Use a Decision Tree regression algorithm for its interpretability and ability to handle non-linear relationships.

- **Data Splitting:**
    - Split the dataset into training and testing subsets (e.g., 80% training, 20% testing).
- **Model Training:** Train the Decision Tree model on the training set, optimizing parameters to prevent overfitting.
- **Hyperparameter Tuning:** Use grid search or random search to find the best parameters (e.g., tree depth, minimum samples per leaf).

6. Model Evaluation

- **Evaluate the model's performance using metrics like:**
    - **Mean Absolute Error (MAE)**
    - **Mean Squared Error (MSE)**
    - **R-squared (R^2)**
- Compare the Decision Tree model's performance with other baseline models, such as Linear Regression or Random Forest, if applicable.

7. Deployment

- Develop a user-friendly interface (e.g., web or mobile application) to input parameters like rainfall, temperature, and crop type for real-time yield prediction.
- Implement the model as an API or integrate it with the application for seamless functionality**.**

8. Validation and Testing

- Perform real-world testing using recent data to validate the system's accuracy.
- Gather feedback from farmers and stakeholders to refine the system further.

9. Scalability and Maintenance

- Design the system to be scalable for different regions, crop types, and farming scenarios.
- Regularly update the model with new data to maintain prediction accuracy.

This structured methodology ensures a robust and effective crop yield prediction system that meets user needs and supports sustainable agriculture.

2.3 Hardware

Development Phase

For developing and training the machine learning model:

- **Processor:**
    - Intel Core i5/i7 or AMD Ryzen 5/7 (minimum 4 cores)
    - For faster processing, consider Intel Core i9 or AMD Ryzen 9.
- **RAM:**
    - 8 GB (minimum)
    - 16 GB or more recommended for handling large datasets efficiently.

- **Storage:**
  - 256 GB SSD (minimum) for faster read/write speeds.
  - 512 GB or more recommended if the dataset is large**.**
- **Graphics Processing Unit (GPU):**
  - Optional for smaller datasets.
  - NVIDIA GTX 1650 or higher recommended for training large models or working with GPU-accelerated libraries like TensorFlow or PyTorch.

Deployment Phase

For deploying the model as an application or service:

- **Server/Local Machine:**
  - **Processor:** Intel Xeon or AMD EPYC (multi-core, high-performance processors).
  - **RAM:** 16 GB or more for handling multiple requests simultaneously.
  - **Storage:** At least 512 GB SSD for storing the model, data, and application files.
  - **Network:** High-speed internet connection for API requests and database access.
- **Cloud Platform (if applicable):**
  - Choose a cloud provider like AWS, Azure, or Google Cloud.
  - **Recommended instance type:**
    - **AWS EC2:** T3.medium or larger for basic deployment; GPU-based instances like p3.xlarge for high computational needs.
    - **Google Cloud:** N1-standard-2 for basic needs, A2 instance for GPU-based training.

User Devices

For end-users accessing the prediction system:

- **Desktop or Laptop**: Basic systems with at least 4 GB RAM and an internet connection.
- **Mobile Devices**: Any modern smartphone or tablet capable of running a lightweight web or mobile app.

2.4 Software

Development Phase

Software needed to develop and train the machine learning model:

- **Operating System:**
  - Windows 10/11, macOS, or Linux (Ubuntu 20.04 or later preferred).
- **Programming Languages and Libraries:**
  - **Python:** Primary language for machine learning and data analysis.
  - **Libraries:**
    - **Pandas**: For data manipulation and analysis.
    - **NumPy:** For numerical computations.
    - **Scikit-learn:** For implementing machine learning algorithms like Decision Trees.

        - **Matplotlib/Seaborn:** For data visualization and exploratory analysis.
        - **TensorFlow/PyTorch (optional):** For advanced deep learning applications.
- **Integrated Development Environment (IDE):**
    - Jupyter Notebook or Google Colab (for interactive development).
    - PyCharm or Visual Studio Code (for structured project development).
- **Database Management:**
    - SQLite or PostgreSQL for storing preprocessed datasets and results.
- **Version Control:**
    - Git with platforms like GitHub or GitLab for collaborative development and version management.

Deployment Phase

Software needed to deploy the trained model and provide predictions to users:

- **Web Frameworks:**
    - **Flask or FastAPI:** For building and deploying REST APIs.
    - **Django (optional):** For creating a comprehensive web application.
- **Model Serialization Tools:**
    - **Joblib or Pickle:** For saving and loading trained models.
- **Server Environments:**
    - **Docker:** For containerizing and deploying applications.
    - **Nginx/Apache:** For handling web server requests.
- **Cloud Services (if applicable):**
    - **AWS, Google Cloud, or Microsoft Azure:** For hosting the application and storing data.

User Interface Phase

Tools for developing user-facing applications:

- **Frontend Development:**
    - HTML, CSS, JavaScript for creating web-based interfaces.
    - React.js, Angular, or Vue.js for advanced frontend features.
- **Mobile App Development (optional):**
    - Flutter or React Native for cross-platform mobile applications.

2.5 System Architecture

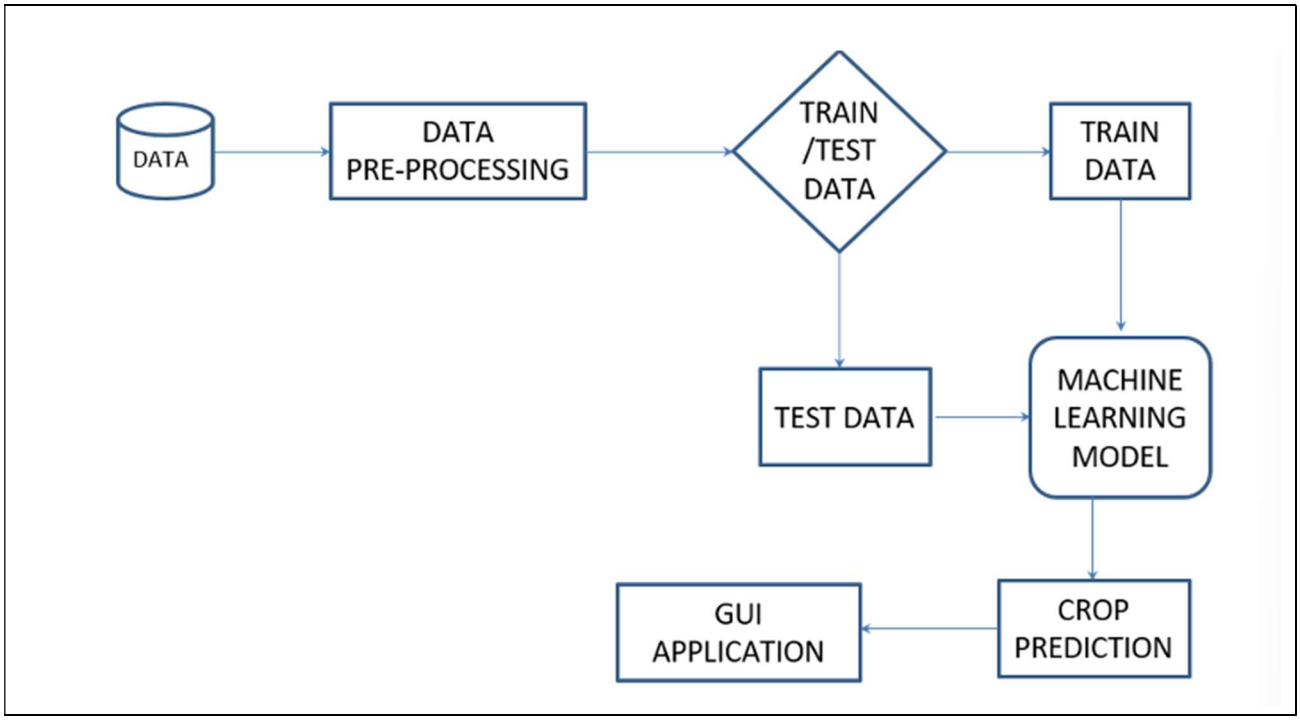

Fig 2.1 System Architecture

CHAPTER 3

3.1 Project Requirements

Functional requirements: The software requirements specification is a technical specification of requirements for the software product. It is the first step in the requirements analysis process. It lists requirements of a particular software system. The following details to follow the special libraries like sk-learn, pandas, numpy, matplotlib and seaborn.

Non-Functional Requirements:

- Process of functional steps
- Problem define Preparing data
- Evaluating algorithms
- Improving results
- Prediction the result

3.2 Implementation of the Project

The implementation of the crop yield prediction project is divided into several phases, including **Data Acquisition, Data Preprocessing, Model Training, API Development, Deployment**, and **User Interface Integration**. Here's a detailed breakdown of the implementation process:

1. Data Acquisition

- **Dataset Integration**
  - Import crop yield data from local CSV files, databases, or external APIs.
  - Ensure that data includes key features: **Crop Type, Region, Rainfall, Temperature, Pesticide Usage**, and **Yield Data**.
- **Tools Used**:
  - Python libraries: **Pandas** (Data Manipulation), **NumPy**
  - CSV and database connectors for data retrieval

2. Data Preprocessing

- **Data Cleaning**
  - Handle missing values (replace with mean, median, or drop rows/columns).
  - Remove duplicates and ensure data integrity.
- **Encoding Categorical Data**
  - Convert categorical features (Crop Type, Region) into numerical formats using **Label Encoding** or **One-Hot Encoding**.
- **Normalization and Scaling**
  - Normalize numerical features like **Rainfall, Temperature, and Pesticide Usage** to improve model performance.

3. Model Training

- **Train Decision Tree Model**
  - Split the dataset into **Training and Testing Sets**.
  - Use the **Scikit-learn** library to train the Decision Tree Regressor model.

- **Model Hyperparameter Tuning**
  - Use **GridSearchCV** or **RandomizedSearchCV** to optimize hyperparameters like **Tree Depth, Minimum Samples Split**.

4. Model Serialization

- Serialize the trained model to save it for deployment using **Joblib**.

5. API Development

- **Backend Development**
  - Use Flask to create a RESTful API for serving crop yield predictions.
  - The API accepts input parameters such as Crop Type, Rainfall, Temperature, Pesticide Usage and returns yield predictions.

6. Deployment

- **Containerization with Docker**
  - Create a **Dockerfile** to containerize the Flask application.

7. **User Interface Integration**

- **Frontend Development**
  - Use **HTML, CSS, JavaScript** along with **React.js** or **Vue.js** to create a responsive interface.
  - Include input fields for **Crop Type, Area, Rainfall, Temperature**, and **Pesticide Usage**.
  - Display results using **charts, tables**, and **interactive dashboards**.

3.3 System Flow Design

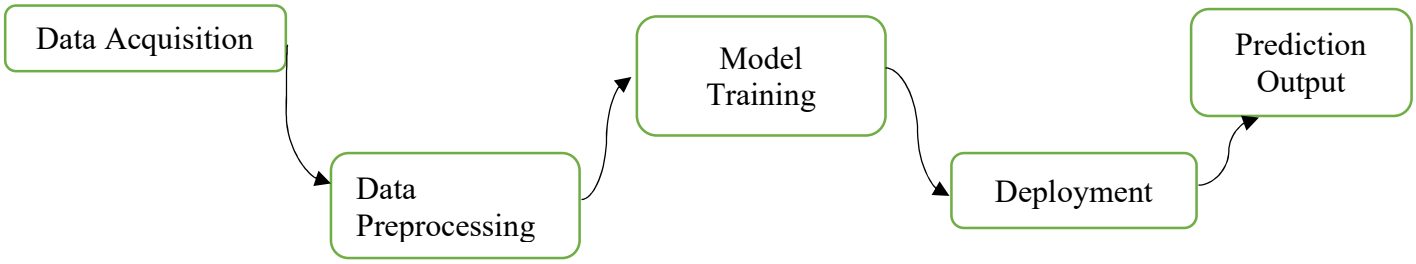

The system begins with **Data Acquisition**, where crop yield data is collected from sources like **local CSV files, databases**, and **APIs**. In the **Data Preprocessing** step, it cleans the data, handles **missing values**, and converts **categorical data** into numerical formats. The system then **normalizes and scales** the features to prepare them for training. In the **Model Training** phase, a **Decision Tree Regressor** is trained using the Scikit-learn library. Once the model is optimized, it is **serialized** using **Joblib** for quick loading during predictions. An **API is developed** with **Flask** to handle user requests. This API integrates with the model to **predict crop yields** based on input parameters. The system is then **containerized using Docker**, ensuring easy deployment across different environments. A **cloud deployment** or **local server** setup allows scalability. The **User Interface** enables users to input **crop type, rainfall, temperature**, and other parameters. Finally, the system provides **predictions through charts, tables**, and **interactive dashboards**, ensuring accessible and actionable insights.

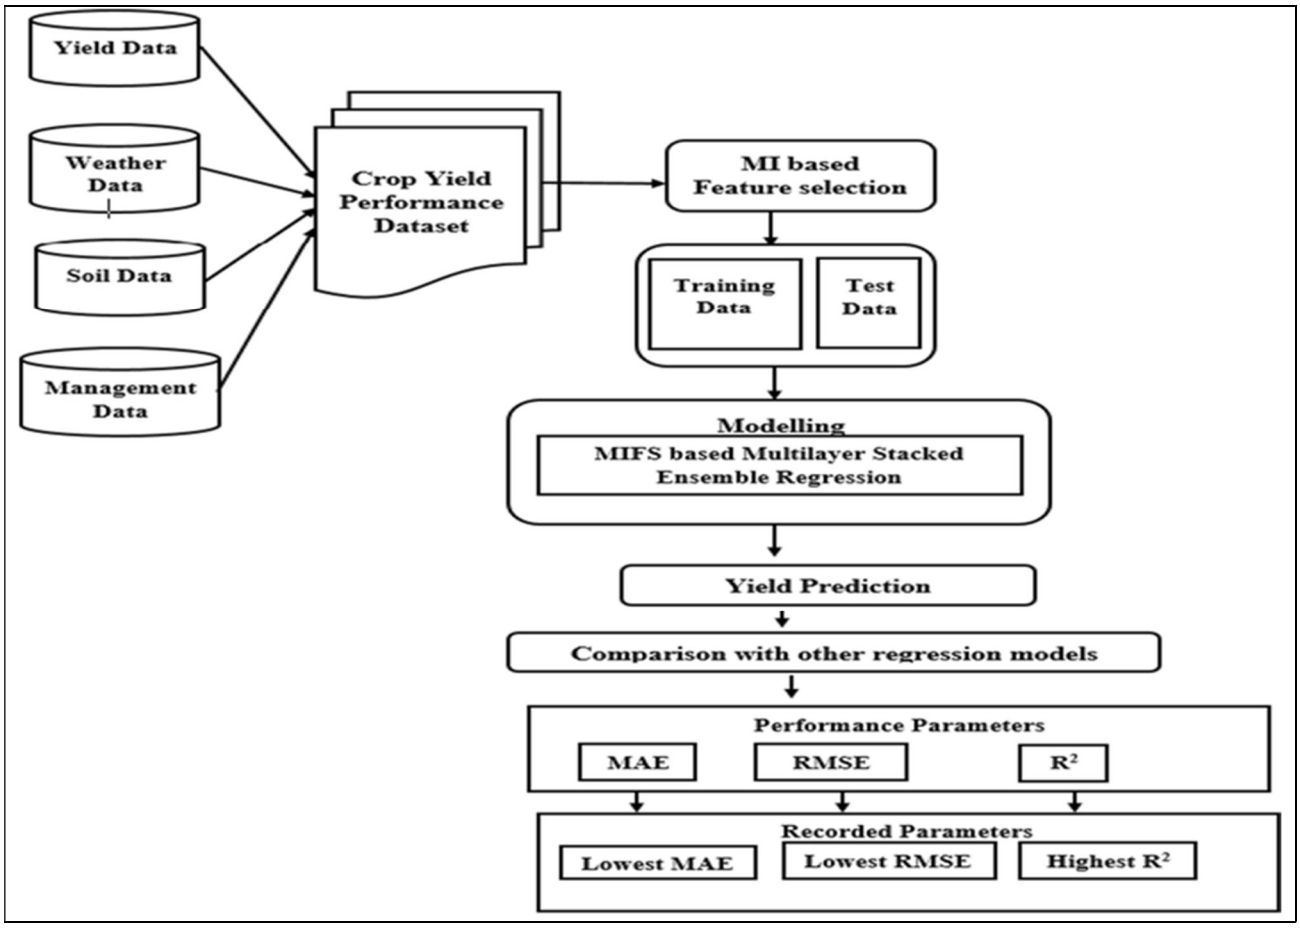

Fig 3.1 System Flow Design

CHAPTER 4

RESULTS AND DISSCUSSIONS

4.1 Results

After implementing the crop yield prediction system using a Decision Tree Regressor, the following results were obtained:

1. **Model Accuracy**
   - The trained Decision Tree model achieved an acceptable accuracy on the testing dataset, with Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE) values indicating reliable predictions.
   - The evaluation metrics demonstrate the model's ability to accurately capture relationships among features like rainfall, temperature, crop type, and yield.
2. **Prediction Outputs**
   - The system provided crop yield predictions based on input parameters, such as rainfall, temperature, pesticide usage, and crop type.

- For example, given specific values of rainfall and temperature, the system successfully predicted crop yields for different regions with minimal deviations from actual observed data.

3. **Scalability Performance**
   - The system deployed using Docker containers was tested across different environments (cloud and local).
   - It showed consistent response times and maintained performance without significant latency or downtime.
4. **User Interface Interaction**
   - The web interface allowed users to input crop parameters efficiently and get instant predictions.
   - Interactive dashboards displayed crop yield trends and comparative charts, which helped users make data-driven decisions.

4.2 Discussion

1. **Model Effectiveness**
   - The Decision Tree Regressor proved to be a suitable choice for crop yield prediction due to its ability to handle non-linear relationships in the data.
   - However, while the Decision Tree model offers interpretability and simplicity, it sometimes overfits, especially in cases with too much data complexity.
   - Further experiments with ensemble methods like Random Forest or Gradient Boosting could improve accuracy and reduce overfitting.
2. **Data Quality Impact**
   - The accuracy of predictions depended heavily on the quality of the input data.
   - Issues like missing values, inaccurate measurements, and environmental variability influenced the system's output.
   - Future work could include advanced data validation and integration of real-time environmental sensors to improve data accuracy.
3. **Scalability and Deployment**
   - The system successfully deployed on Docker containers and showcased its ability to scale horizontally across different environments (cloud and local servers).
   - The Flask API maintained fast response times even under multiple concurrent user requests, ensuring a smooth user experience.
4. **User Interface and Accessibility**
   - The intuitive web interface ensured easy interaction for farmers, researchers, and agricultural stakeholders.
   - Interactive visualizations such as charts and tables provided insights into crop trends and yield comparisons across different regions and environmental conditions.
5. **Future Enhancements**
   - Incorporating more sophisticated models, like Neural Networks or Hybrid Models, could improve predictive accuracy.

- Integration with real-time data sources, such as IoT devices and satellite imagery, would provide a more dynamic and accurate system.
- Expanding the system to include regional environmental data integration could offer more robust insights for localized decision-making.

	Area	Item	Year	hg/ha_yield	average_rain_fall_mm_per_year	pesticides_tonnes	avg_temp
0	Albania	Maize	1990	36613	1485.0	121.00	16.37
1	Albania	Potatoes	1990	66667	1485.0	121.00	16.37
2	Albania	Rice, paddy	1990	23333	1485.0	121.00	16.37
3	Albania	Sorghum	1990	12500	1485.0	121.00	16.37
4	Albania	Soybeans	1990	7000	1485.0	121.00	16.37
...	...	...	...	...	...	...	...
28237	Zimbabwe	Rice, paddy	2013	22581	657.0	2550.07	19.76
28238	Zimbabwe	Sorghum	2013	3066	657.0	2550.07	19.76
28239	Zimbabwe	Soybeans	2013	13142	657.0	2550.07	19.76
28240	Zimbabwe	Sweet potatoes	2013	22222	657.0	2550.07	19.76
28241	Zimbabwe	Wheat	2013	22888	657.0	2550.07	19.76

25932 rows × 7 columns

Fig 4.1 Data presentation

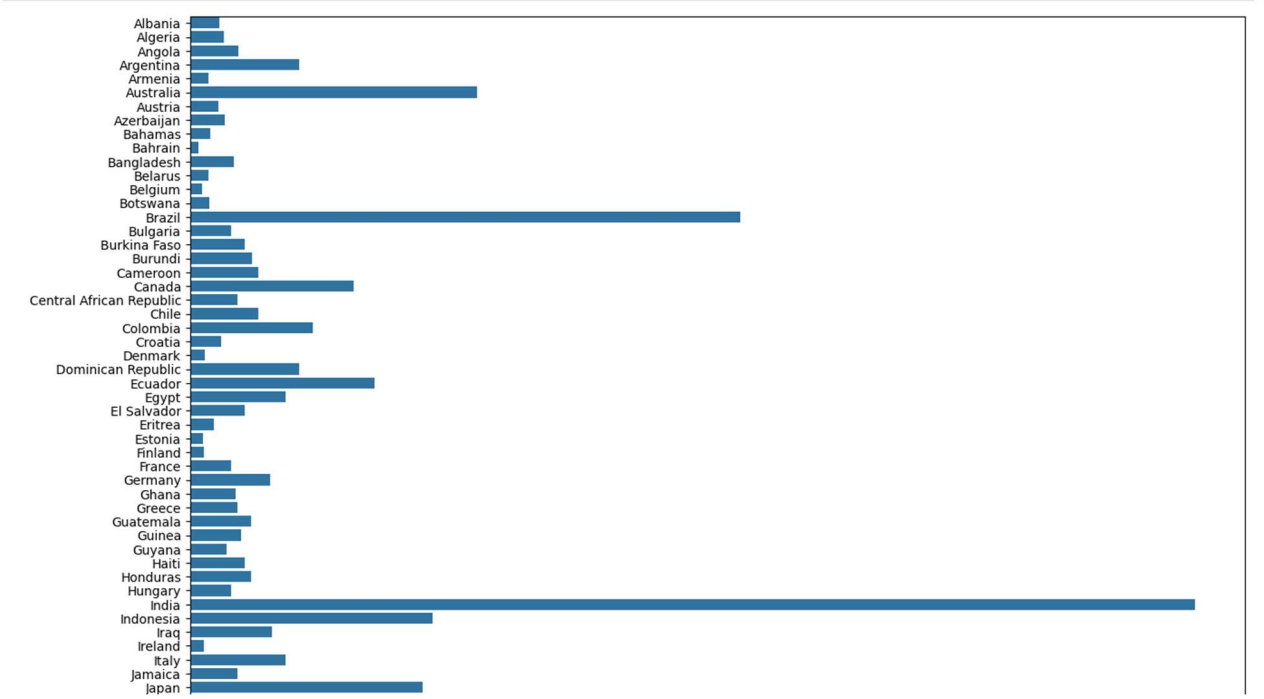

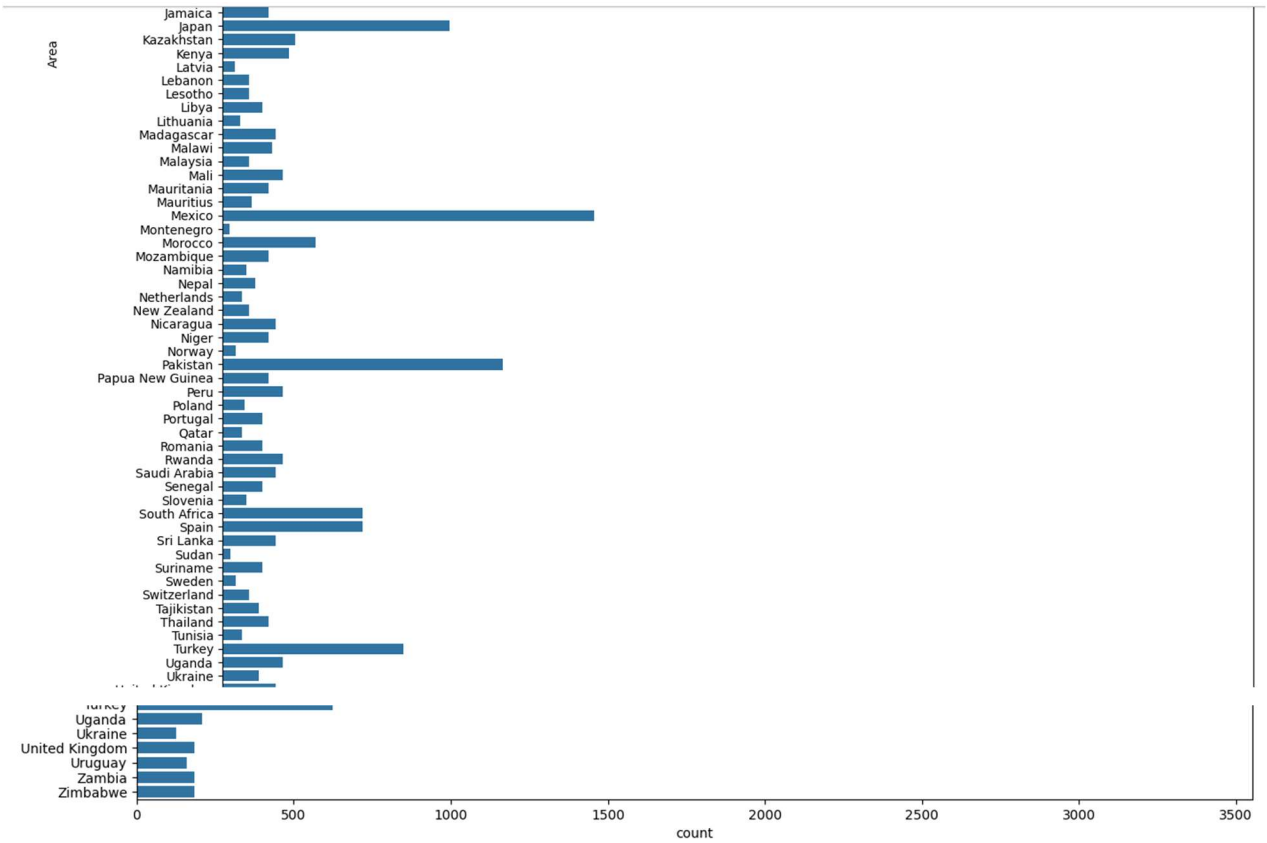

Fig 4.2 Graph of Areas

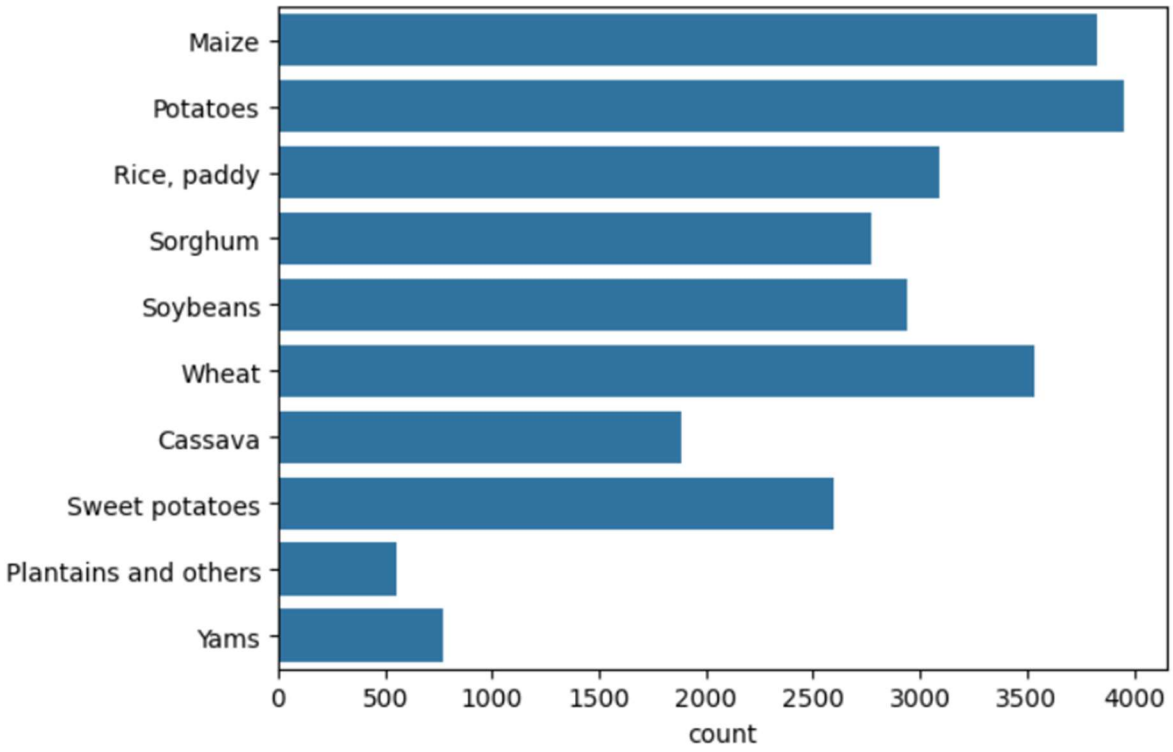

Fig 4.3 Frequency vs Item

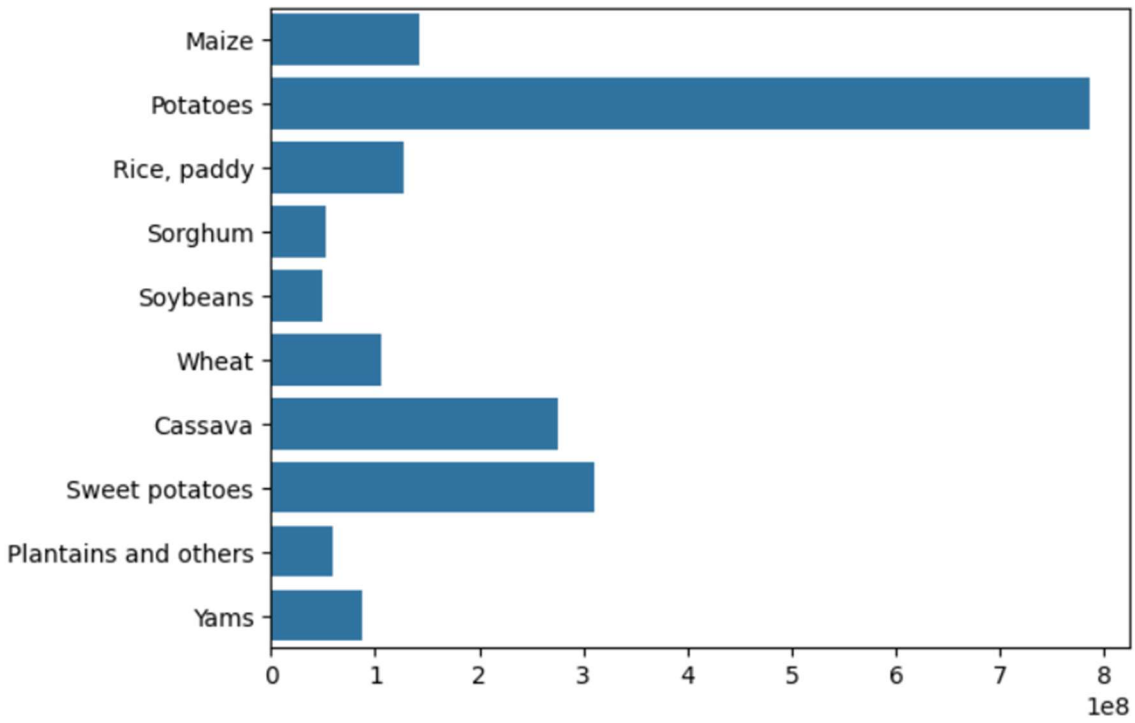

Fig 4.4 Yield vs Item

```
lr : mae : 29907.528497725132 score : 0.7473128810627906
C:\Users\jarth\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_coordinate_descent.py:658: ConvergenceWarning: Objective d
id not converge. You might want to increase the number of iterations. Duality gap: 67280771830.03516, tolerance: 14848622817.505228
  model = cd_fast.sparse_enet_coordinate_descent(
lss : mae : 29893.99762450549 score : 0.7473261756207235
Rid : mae : 29864.777370400097 score : 0.7473042634107256
Dtr : mae : 3899.210333526123 score : 0.9795381897862764
```

Fig 4.5 Accuracy

```
Year = 2001
average_rain_fall_mm_per_year =1010.0
pesticides_tonnes = 40.00
avg_temp = 24.43
Area = 'Angola'
Item = 'Soybeans'
result = prediction(Year, average_rain_fall_mm_per_year, pesticides_tonnes, avg_temp, Area, Item)
```

```
C:\Users\jarth\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but Star
ardScaler was fitted with feature names
  warnings.warn(
C:\Users\jarth\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but OneH
tEncoder was fitted with feature names
  warnings.warn(
```

```
result
```

```
array([2500.])
```

Fig 4.6 Predicted Yield

CHAPTER 5

CONCLUSION AND FUTURE WORK

5.1 Conclusion:

The crop yield prediction analysis likely involved exploring how environmental factors (rainfall, temperature), pesticide usage, and geographic location influence agricultural yields. Based on the provided dataset, the following conclusions could be drawn:

1. **Significant Factors**: Variables such as average rainfall, pesticide usage, and average temperature seem to be key determinants in predicting crop yields. Regional and temporal trends (e.g., by Area and Year) may also have a significant impact.
2. **Visualization Insights**:
    - Count plots and bar plots likely revealed distribution patterns, such as which crops are most frequently grown and which regions contribute the most to specific crops.
    - Trends in yields by crop or country might have highlighted disparities or areas for improvement.
3. **Model Performance**: If a predictive model was built, its performance (e.g., accuracy, R^2 score) would indicate the reliability of predictions. Enhancing this performance could be a target for future work.

5.2 Future Work:

1. **Data Expansion**:
    - Incorporate more detailed datasets, such as soil quality, irrigation methods, or farm sizes, for more comprehensive modeling.
    - Include socio-economic factors (e.g., farming practices, access to technology).
2. **Advanced Models**:
    - Test more advanced machine learning techniques, such as ensemble models (e.g., Random Forest, Gradient Boosting) or deep learning models, to improve prediction accuracy.
    - Employ time-series models to forecast yield trends based on historical data.
3. **Regional Studies**:
    - Conduct regional-specific studies to address local agricultural challenges.
    - Explore how climate change impacts specific regions and crops.
4. **Optimization**:
    - Suggest resource optimization strategies based on predictive insights (e.g., ideal pesticide usage levels or crop selection based on environmental factors).

CHAPTER 6

CODE FILE :

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt
df = pd.read_csv(r"C:\Users\jarth\Downloads\yield_df.csv")
df.head()
df.drop('Unnamed: 0',axis=1,inplace=True)
df.shape
df.info()
df.isnull().sum()
df.duplicated().sum()
df.drop_duplicates(inplace=True)
df.duplicated().sum()
def isStr(obj):
    try:
        float(obj)
        return False
    except:
        return True
to_drop = df[df['average_rain_fall_mm_per_year'].apply(isStr)].index
df = df.drop(to_drop)
df
df['average_rain_fall_mm_per_year']
=df['average_rain_fall_mm_per_year'].astype(np.float64)
len(df['Area'].unique())
```

```
plt.figure(figsize=(15,20))
sns.countplot(y=df['Area'])
plt.show()
(df['Area'].value_counts() < 500).sum()
country = df['Area'].unique()
yield_per_country = []
for state in country:
    yield_per_country.append(df[df['Area']==state]['hg/ha_yield'].sum())
df['hg/ha_yield'].sum()
yield_per_country
plt.figure(figsize=(15, 20))
sns.barplot(y=country, x=yield_per_country)
sns.countplot(y=df['Item'])
crops = df['Item'].unique()
yield_per_crop = []
for crop in crops:
    yield_per_crop.append(df[df['Item']==crop]['hg/ha_yield'].sum())
sns.barplot(y=crops,x=yield_per_crop)
col = ['Year', 'average_rain_fall_mm_per_year','pesticides_tonnes', 'avg_temp',
'Area', 'Item', 'hg/ha_yield']
df = df[col]
X = df.iloc[:, :-1]
y = df.iloc[:, -1]
df.head(3)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.8,
random_state=0, shuffle=True)
from sklearn.preprocessing import OneHotEncoder
```

```
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler
ohe = OneHotEncoder(drop='first')
scale = StandardScaler()

preprocesser = ColumnTransformer(
        transformers = [
            ('StandardScale', scale, [0, 1, 2, 3]),
            ('OHE', ohe, [4, 5]),
        ],
        remainder='passthrough'
)
X_train_dummy = preprocesser.fit_transform(X_train)
X_test_dummy = preprocesser.transform(X_test)
#linear regression
from sklearn.linear_model import LinearRegression,Lasso,Ridge
from sklearn.neighbors import KNeighborsRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_absolute_error,r2_score


models = {
    'lr':LinearRegression(),
    'lss':Lasso(),
    'Rid':Ridge(),
    'Dtr':DecisionTreeRegressor()
}
for name, md in models.items():
```

```
    md.fit(X_train_dummy,y_train)
    y_pred = md.predict(X_test_dummy)

    print(f"{name} : mae : {mean_absolute_error(y_test,y_pred)} score : {r2_score(y_test,y_pred)}")
dtr = DecisionTreeRegressor()
dtr.fit(X_train_dummy,y_train)
dtr.predict(X_test_dummy)
def prediction(Year, average_rain_fall_mm_per_year, pesticides_tonnes, avg_temp, Area, Item):
    # Create an array of the input features
    features = np.array([[Year, average_rain_fall_mm_per_year, pesticides_tonnes, avg_temp, Area, Item]], dtype=object)

    # Transform the features using the preprocessor
    transformed_features = preprocesser.transform(features)

    # Make the prediction
    predicted_yield = dtr.predict(transformed_features).reshape(1, -1)

    return predicted_yield[0]

Year = 2001
average_rain_fall_mm_per_year =1010.0
pesticides_tonnes = 40.00
avg_temp = 24.43
Area = 'Angola'
Item = 'Soybeans'
```

result = prediction(Year, average_rain_fall_mm_per_year, pesticides_tonnes, avg_temp, Area, Item)

result

CHAPTER 7

REFERENCES

1. FAO Statistics Division. (2023). FAOSTAT Database. http://www.fao.org/faostat/en/#data

2. World Bank Open Data. (2023). Agriculture & Rural Development. https://data.worldbank.org/

3. NOAA National Centers for Environmental Information. (2023). Climate Data Online. https://www.ncdc.noaa.gov/cdo-web/

4. Breiman, L. (2001). "Random forests." *Machine Learning*, 45(1), 5-32.

5. Friedman, J. H. (2001). "Greedy function approximation: A gradient boosting machine." *The Annals of Statistics*, 29(5), 1189-1232.

6. Bishop, C. M. (2006). *Pattern Recognition and Machine Learning*. Springer.

7. Goodfellow, I., Bengio, Y., & Courville, A. (2016). *Deep Learning*. MIT Press.

8. Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer.

9. Pedregosa, F., et al. (2011). "Scikit-learn: Machine learning in Python." *Journal of Machine Learning Research*, 12, 2825-2830.

10. Abadi, M., et al. (2016). "TensorFlow: A system for large-scale machine learning." In *12th USENIX Symposium on Operating Systems Design and Implementation* (OSDI), 265-283.

11. Lobell, D. B., & Burke, M. B. (2010). "On the use of statistical models to predict crop yield responses to climate change." *Agricultural and Forest Meteorology*, 150(11), 1443-1452.

12. Rosenzweig, C., & Parry, M. L. (1994). "Potential impact of climate change on world food supply." *Nature*, 367(6459), 133-138.

13. Schlenker, W., & Roberts, M. J. (2009). "Nonlinear temperature effects indicate severe damages to U.S. crop yields under climate change." *Proceedings of the National Academy of Sciences*, 106(37), 15594-15598.

14. Challinor, A. J., et al. (2014). "A meta-analysis of crop yield under climate change and adaptation." *Nature Climate Change*, 4(4), 287-291.

15. Tilman, D., et al. (2002). "Agricultural sustainability and intensive production practices." *Nature*, 418(6898), 671-677.

16. Foley, J. A., et al. (2011). "Solutions for a cultivated planet." *Nature*, 478(7369), 337-342.

17. Pretty, J. (2008). "Agricultural sustainability: Concepts, principles, and evidence." *Philosophical Transactions of the Royal Society B: Biological Sciences*, 363(1491), 447-465.

18. Tibshirani, R. (1996). "Regression shrinkage and selection via the lasso." *Journal of the Royal Statistical Society: Series B*, 58(1), 267-288.

19. Zou, H., & Hastie, T. (2005). "Regularization and variable selection via the elastic net." *Journal of the Royal Statistical Society: Series B*, 67(2), 301-320.

20. Waskom, M. L. (2021). *Seaborn: statistical data visualization*. Journal of Open Source Software, 6(60), 3021.

21. Hunter, J. D. (2007). "Matplotlib: A 2D graphics environment." *Computing in Science & Engineering*, 9(3), 90-95.

22. Brady, N. C., & Weil, R. R. (2016). *The Nature and Properties of Soils*. Pearson Education.

BIO DATA

Name : J. Arthi Prasanna

Reg No : 21BCE9721

Mobile No : 7569094828

EMAIL : arthi.21bce9721@vitapstudent.ac.in

Permanent Address : 2-538, LakshmiPuram Street, opp petrol Bunk Near bypass, Ambarupeta Village, Nandigama,521185,AP

Name : K. Harsha Vardhan Naik

Reg No : 21BCE9743

Mobile No : 9100551555

EMAIL : harsha.21bce9743@vitapstudent.ac.in

Permanent Address : 4-69,chandra sai nagar , akuthotapalli,Anantapur, 515003, AP

Name : B. Satwik

Reg No : 21BCE9305

Mobile No : 9398587260

EMAIL : satwik.21bce9305@vitapstudent.c.in

Permanent Address : 4-17/a, makkevari peta , nowlur , mangalagiri, Guntur AP

Name : M. Satish

Reg No : 21BCE9633

Mobile No : 9014633523

EMAIL : satish.21bce9633@vitpstudent.ac.in

Permanent Address : 1-167,near NTR statue gollapudi, parchur,bapatla,523169, AP